

```
[15] import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
import xgboost as xgb
from sklearn.metrics import r2_score, recall_score, f1_score,
roc_curve, roc_auc_score
from sklearn.ensemble import GradientBoostingClassifier
import pandas as pd
from sklearn.metrics import accuracy_score
from xgboost import XGBClassifier
from sklearn import metrics
```



```
[2] def plot_roc_curve(y_train, preds_train, y_test, preds_test):
    plt.plot(metrics.roc_curve(y_train, preds_train)[0],
metrics.roc_curve(y_train, preds_train)[1],
            color = 'red', label='Train ROC Curve (area =
%0.5f)' % roc_auc_score(y_train, preds_train))
    plt.plot(metrics.roc_curve(y_test, preds_test)
[0],metrics.roc_curve(y_test, preds_test)[1],
            color = 'blue', label='Test ROC Curve (area =
%0.5f)' % roc_auc_score(y_test, preds_test))
    plt.plot([0, 2], [0, 2], color='black', linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.0])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('AUC')
    plt.legend()
    plt.show()
    sns.set(style='white', rc={'figure.figsize':(10,10)})
```

```
[3] def important_stats(y_true, y_pred_proba, summary):
    print("-----")
    y_pred_label = pd.Series(y_pred_proba)
    y_pred_label = y_pred_label.map(lambda x: 1 if x > 0.5 else
0)
    print(summary)
    recall = recall_score(y_true, y_pred_label)
    print('recall:', recall)
    f1_stat = f1_score(y_true, y_pred_label)
    print('f1_score:', f1_stat)
    accuracyScore= accuracy_score(y_true, y_pred_label)
    print('accuracy_score:', accuracyScore)
    fpr, tpr, thresholds = metrics.roc_curve(y_true,
y_pred_proba)
    auc = metrics.auc(fpr, tpr)
    print('AUC:', auc)
```

```
matrix = pd.crosstab(y_true, y_pred_label, rownames=['True'],
colnames=['Predicted'], margins=True)
print(matrix)
print("-----")
```

```
[4] df =
pd.read_csv(f'{os.getcwd()}/data_clean/cleaned_match_data.csv')
```

```
[5] df.head()
```

	Unnamed: 0	match_id	duration	result	top_towers	mid_tower
0	0	0	2375	1	1	3
1	1	0	2375	0	-1	-3
2	2	0	2375	1	1	3
3	3	0	2375	0	-1	-3
4	4	1	2582	0	-2	-2

5 rows × 54 columns

```
[6] df.shape
```

(97342, 54)

```
[83] df.dtypes
```

```
Unnamed: 0          int64
match_id           int64
duration           int64
result             int64
top_towers         int64
mid_towers         int64
bottom_towers      int64
ancient_status    int64
top_barracks       int64
mid_barracks       int64
bottom_barracks    int64
gold_total         int64
gold_max           int64
gold_min           int64
gold_std           float64
gold_spent_avg     float64
```

```

gold_spent_max      int64
gold_spent_min      int64
gold_spent_std      float64
kills_total         int64
deaths_total        int64
deaths_max          int64
deaths_min          int64
deaths_std          float64
assists_avg         float64
assists_max         int64
assists_min         int64
assists_std         float64
denies_avg          float64
denies_max          int64
denies_min          int64
denies_std          float64
last_hits_avg       float64
last_hits_max       int64
last_hits_min       int64
last_hits_std       float64
hero_damage_total   int64
hero_damage_max     int64
hero_damage_min     int64
hero_damage_std     float64
tower_damage_total  int64
tower_damage_max    int64
tower_damage_min    int64
tower_damage_std    float64
level_total         int64
level_max           int64
level_min           int64
level_std           float64
gold_buyback_avg    float64
gold_buyback_max    float64
gold_buyback_min    float64
gold_buyback_std    float64
teamfight_loss      int64
has_negative_chat    bool
dtype: object

```

```

[7] df.head()
df = df.drop(columns = ['match_id', 'Unnamed: 0'])

```

```

[8] df.corr()

```

	duration	result	top_towers	mid_towers
duration	1.000000	0.000000	0.000000	0.000000
result	0.000000	1.000000	0.780384	0.869120
top_towers	0.000000	0.780384	1.000000	0.807989

	duration	result	top_towers	mid_towers	
mid_towers	0.000000	0.869120	0.807989	1.000000	
bottom_towers	0.000000	0.825153	0.792817	0.856402	
ancient_status	0.000000	0.986854	0.790515	0.880311	
top_barracks	0.000000	0.809052	0.880293	0.738965	
mid_barracks	0.000000	0.865915	0.752831	0.801983	
bottom_barracks	0.000000	0.942938	0.787292	0.914214	
gold_total	0.230393	0.739365	0.548157	0.613873	
gold_max	0.221347	0.657264	0.493350	0.550434	
gold_min	0.167829	0.546885	0.399394	0.443538	
gold_std	0.188805	0.543037	0.411468	0.459860	
gold_spent_avg	0.785802	0.484111	0.449560	0.489240	
gold_spent_max	0.688970	0.423778	0.394271	0.422551	
gold_spent_min	0.656757	0.435685	0.407193	0.449461	
gold_spent_std	0.484055	0.281615	0.260556	0.271514	
kills_total	0.573466	0.533060	0.476218	0.559674	
deaths_total	0.580572	-0.524347	-0.466332	-0.548470	
deaths_max	0.556368	-0.439920	-0.396471	-0.470060	
deaths_min	0.474530	-0.509340	-0.445814	-0.519711	
deaths_std	0.315636	-0.114377	-0.115589	-0.146337	
assists_avg	0.588018	0.447070	0.355174	0.441683	
assists_max	0.556651	0.450628	0.367368	0.450488	
assists_min	0.526481	0.365420	0.278232	0.355823	
assists_std	0.350211	0.343699	0.299277	0.352546	
denies_avg	0.089869	0.116800	0.133224	0.145273	
denies_max	0.059878	0.091193	0.105878	0.114093	
denies_min	0.081750	0.059385	0.061330	0.069323	
denies_std	0.046099	0.083094	0.098017	0.105110	

	duration	result	top_towers	mid_towers	
last_hits_avg	0.848630	0.165491	0.177155	0.177493	
last_hits_max	0.736927	0.175365	0.184410	0.178215	
last_hits_min	0.602318	0.068833	0.072592	0.087823	
last_hits_std	0.633463	0.174293	0.181553	0.171999	
hero_damage_total	0.716266	0.366720	0.337051	0.393060	
hero_damage_max	0.584585	0.367620	0.333298	0.389706	
hero_damage_min	0.493316	0.197360	0.189298	0.219091	
hero_damage_std	0.428317	0.320831	0.285614	0.335256	
tower_damage_total	0.179902	0.883605	0.815497	0.858416	
tower_damage_max	0.192582	0.785854	0.731209	0.757441	
tower_damage_min	0.006648	0.613200	0.560517	0.619049	
tower_damage_std	0.199974	0.747728	0.696752	0.716737	
level_total	0.867513	0.320456	0.278445	0.324326	
level_max	0.776655	0.318169	0.296044	0.338097	
level_min	0.780949	0.293050	0.246091	0.292922	
level_std	0.063222	0.059137	0.088417	0.083004	
gold_buyback_avg	-0.585443	0.189751	0.141134	0.142127	
gold_buyback_max	-0.213388	0.061799	0.036464	0.024962	
gold_buyback_min	-0.602811	0.161717	0.121168	0.131745	
gold_buyback_std	0.582065	-0.154209	-0.118323	-0.133527	
teamfight_loss	0.242851	-0.565136	-0.517809	-0.609866	
has_negative_chat	0.047286	0.000000	0.000000	0.000000	

52 rows × 52 columns

```
[9] #Drop highly correlated features
corr_matrix = df.corr().abs()
upper =
corr_matrix.where(np.triu(np.ones(corr_matrix.shape),k=1).astype(
np.bool))
```

```
to_drop = [column for column in upper.columns if
any(upper[column]>0.6)]
df.drop(to_drop,axis=1,inplace=True)
df.head()
```

	duration	result	deaths_total	denies_avg	denies_min	level_
0	2375	1	17	6.0	1	3.0332
1	2375	0	52	7.6	0	2.6382
2	2375	1	17	6.0	1	3.0332
3	2375	0	52	7.6	0	2.6382
4	2582	0	53	5.4	0	4.4091

```
[10] x_train, x_test, y_train, y_test = train_test_split(
      df.drop(columns = ['result','duration']),
      df['result'],
      test_size=0.2,
      random_state=1
    )
```

```
[11] x_train, x_val, y_train, y_val = train_test_split(x_train,
      y_train, test_size = 0.2, random_state = 1)
```

XGBoost

```
[12] xgb = XGBClassifier(random_state=0, n_jobs=-1, learning_rate=0.1,
      n_estimators=100, max_depth=3)
model = xgb.fit(x_train, y_train)
y_pred_test = xgb.predict_proba(x_test)[:, 1]
y_pred_train = xgb.predict_proba(x_train)[:, 1]
important_stats(y_train, y_pred_train, "train result summary: ")
important_stats(y_test, y_pred_test, "test result summary: ")
```

```
-----
train result summary:
recall: 0.694505071254333
f1_score: 0.7549368501849137
accuracy_score: 0.7745031943240553
AUC: 0.853348040516045
Predicted      0      1    All
```

True			
0	11525	8389	19914
1	11625	8272	19897
All	23150	16661	39811

test result summary:

recall: 0.6991903248949473

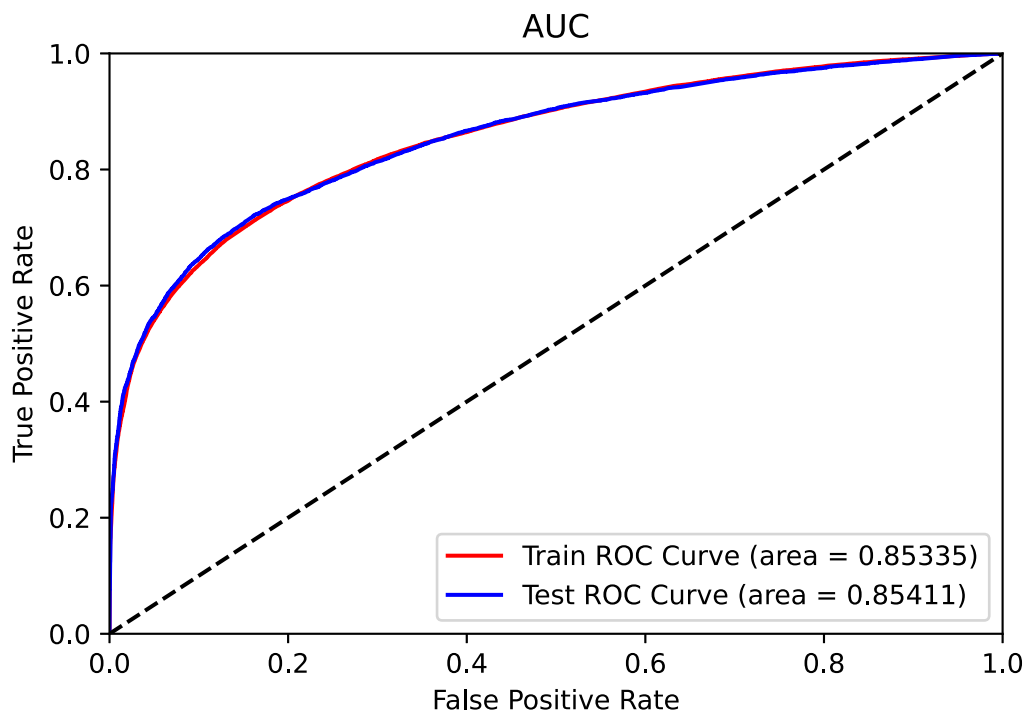
f1_score: 0.7596881959910914

accuracy_score: 0.7783142431557861

AUC: 0.8541094150037003

Predicted	0	1	All
True			
0	1146	849	1995
1	1151	864	2015
All	2297	1713	4010

```
[13] import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
plot_roc_curve(y_train, y_pred_train, y_test, y_pred_test)
```



GBDT

```
[64] from sklearn.ensemble import GradientBoostingClassifier
```

```

gbdt = GradientBoostingClassifier(random_state=0,
n_estimators=10, max_depth=10)
gbdt = gbdt.fit(x_train, y_train)
y_pred_test = gbdt.predict_proba(x_test)[:, 1]
y_pred_train = gbdt.predict_proba(x_train)[:, 1]
important_stats(y_train, y_pred_train, "train result summary: ")
important_stats(y_test, y_pred_test, "test result summary: ")

```

train result summary:

recall: 0.7515727307741688

f1_score: 0.8036655054656531

accuracy_score: 0.816350444637067

AUC: 0.8934285595601786

Predicted	0	1	All
-----------	---	---	-----

True			
------	--	--	--

0	11213	8701	19914
---	-------	------	-------

1	11324	8573	19897
---	-------	------	-------

All	22537	17274	39811
-----	-------	-------	-------

test result summary:

recall: 0.7122066208875679

f1_score: 0.7600765654908396

accuracy_score: 0.7746674200010273

AUC: 0.8496263887085502

Predicted	0	1	All
-----------	---	---	-----

True			
------	--	--	--

0	1109	886	1995
---	------	-----	------


1	1124	891	2015
---	------	-----	------

All	2233	1777	4010
-----	------	------	------

```

[66] plot_roc_curve(y_train, y_pred_train, y_test, y_pred_test)

```

LightGBM

```
[107] import lightgbm as gbm
      from lightgbm import LGBMClassifier
```

```
[108] gbm_clf = gbm.LGBMClassifier(
    boosting_type = 'gbdt',
    #num_leaves = ,
    #max_depth = ,
    learning_rate = 0.1
    #n_estimators =
    #,subsample_for_bin =
    ,objective = 'binary'
    ,metric = 'binary_logloss'
    #,class_weight =
    #,min_split_gain =
    #,min_split_weight =
    #,min_child_weight =
    #,min_child_samples =
    #,subsample =
    #,subsample_freq =
    #,colsample_bytree =
    ,reg_alpha = 5
    ,reg_lambda = 120
    ,importance_type = 'split' #will rank features by # of times
    it is used in model.'gain' for gain
    ,num_iterations = 1000
)
```

```
[109] gbm_clf.fit(
    x_train,
    y_train,
    eval_metric = 'result',
    verbose = True,
    eval_set = [(x_val, y_val)],
    early_stopping_rounds = 20
)
```

```
[1] valid_0's binary_logloss: 0.662348
Training until validation scores don't improve for 20 rounds
[2] valid_0's binary_logloss: 0.637353
[3] valid_0's binary_logloss: 0.615892
[4] valid_0's binary_logloss: 0.597688
[5] valid_0's binary_logloss: 0.582505
[6] valid_0's binary_logloss: 0.569258
[7] valid_0's binary_logloss: 0.55816
[8] valid_0's binary_logloss: 0.548532
[9] valid_0's binary_logloss: 0.539825
[10] valid_0's binary_logloss: 0.532286
[11] valid_0's binary_logloss: 0.52577
[12] valid_0's binary_logloss: 0.519981
[13] valid_0's binary_logloss: 0.515134
[14] valid_0's binary_logloss: 0.510994
[15] valid_0's binary_logloss: 0.507042
```

[16] valid_0's binary_logloss: 0.50359
[17] valid_0's binary_logloss: 0.500632
[18] valid_0's binary_logloss: 0.497876
[19] valid_0's binary_logloss: 0.49563
[20] valid_0's binary_logloss: 0.493652
[21] valid_0's binary_logloss: 0.491747
[22] valid_0's binary_logloss: 0.4902
[23] valid_0's binary_logloss: 0.488692
[24] valid_0's binary_logloss: 0.487449
[25] valid_0's binary_logloss: 0.486344
[26] valid_0's binary_logloss: 0.485169
[27] valid_0's binary_logloss: 0.484033
[28] valid_0's binary_logloss: 0.483117
[29] valid_0's binary_logloss: 0.482207
[30] valid_0's binary_logloss: 0.481406
[31] valid_0's binary_logloss: 0.480741
[32] valid_0's binary_logloss: 0.480101
[33] valid_0's binary_logloss: 0.479612
[34] valid_0's binary_logloss: 0.479081
[35] valid_0's binary_logloss: 0.478507
[36] valid_0's binary_logloss: 0.47803
[37] valid_0's binary_logloss: 0.477629
[38] valid_0's binary_logloss: 0.477236
[39] valid_0's binary_logloss: 0.476853
[40] valid_0's binary_logloss: 0.476455
[41] valid_0's binary_logloss: 0.476112
[42] valid_0's binary_logloss: 0.47588
[43] valid_0's binary_logloss: 0.475598
[44] valid_0's binary_logloss: 0.475332
[45] valid_0's binary_logloss: 0.475145
[46] valid_0's binary_logloss: 0.474895
[47] valid_0's binary_logloss: 0.474689
[48] valid_0's binary_logloss: 0.474463
[49] valid_0's binary_logloss: 0.474345
[50] valid_0's binary_logloss: 0.474202
[51] valid_0's binary_logloss: 0.474086
[52] valid_0's binary_logloss: 0.4739
[53] valid_0's binary_logloss: 0.473702
[54] valid_0's binary_logloss: 0.473563
[55] valid_0's binary_logloss: 0.473427
[56] valid_0's binary_logloss: 0.473301
[57] valid_0's binary_logloss: 0.473238
[58] valid_0's binary_logloss: 0.473103
[59] valid_0's binary_logloss: 0.47303
[60] valid_0's binary_logloss: 0.472866
[61] valid_0's binary_logloss: 0.472765
[62] valid_0's binary_logloss: 0.472737
[63] valid_0's binary_logloss: 0.472584
[64] valid_0's binary_logloss: 0.472511
[65] valid_0's binary_logloss: 0.472458
[66] valid_0's binary_logloss: 0.472435
[67] valid_0's binary_logloss: 0.472354

[68] valid_0's binary_logloss: 0.472355
[69] valid_0's binary_logloss: 0.472297
[70] valid_0's binary_logloss: 0.472185
[71] valid_0's binary_logloss: 0.472137
[72] valid_0's binary_logloss: 0.472095
[73] valid_0's binary_logloss: 0.472048
[74] valid_0's binary_logloss: 0.472025
[75] valid_0's binary_logloss: 0.471999
[76] valid_0's binary_logloss: 0.471925
[77] valid_0's binary_logloss: 0.471845
[78] valid_0's binary_logloss: 0.471815
[79] valid_0's binary_logloss: 0.471743
[80] valid_0's binary_logloss: 0.471733
[81] valid_0's binary_logloss: 0.471701
[82] valid_0's binary_logloss: 0.471631
[83] valid_0's binary_logloss: 0.471544
[84] valid_0's binary_logloss: 0.471519
[85] valid_0's binary_logloss: 0.471492
[86] valid_0's binary_logloss: 0.471465
[87] valid_0's binary_logloss: 0.471421
[88] valid_0's binary_logloss: 0.47136
[89] valid_0's binary_logloss: 0.471333
[90] valid_0's binary_logloss: 0.471315
[91] valid_0's binary_logloss: 0.471279
[92] valid_0's binary_logloss: 0.471268
[93] valid_0's binary_logloss: 0.471213
[94] valid_0's binary_logloss: 0.471181
[95] valid_0's binary_logloss: 0.471163
[96] valid_0's binary_logloss: 0.471141
[97] valid_0's binary_logloss: 0.47112
[98] valid_0's binary_logloss: 0.471043
[99] valid_0's binary_logloss: 0.471022
[100] valid_0's binary_logloss: 0.470953
[101] valid_0's binary_logloss: 0.470941
[102] valid_0's binary_logloss: 0.470935
[103] valid_0's binary_logloss: 0.470899
[104] valid_0's binary_logloss: 0.470889
[105] valid_0's binary_logloss: 0.470889
[106] valid_0's binary_logloss: 0.470875
[107] valid_0's binary_logloss: 0.470829
[108] valid_0's binary_logloss: 0.470829
[109] valid_0's binary_logloss: 0.470807
[110] valid_0's binary_logloss: 0.470772
[111] valid_0's binary_logloss: 0.470746
[112] valid_0's binary_logloss: 0.470724
[113] valid_0's binary_logloss: 0.470666
[114] valid_0's binary_logloss: 0.470613
[115] valid_0's binary_logloss: 0.470623
[116] valid_0's binary_logloss: 0.470581
[117] valid_0's binary_logloss: 0.470596
[118] valid_0's binary_logloss: 0.470589
[119] valid_0's binary_logloss: 0.470559

[120] valid_0's binary_logloss: 0.470511
[121] valid_0's binary_logloss: 0.470501
[122] valid_0's binary_logloss: 0.470474
[123] valid_0's binary_logloss: 0.470504
[124] valid_0's binary_logloss: 0.470503
[125] valid_0's binary_logloss: 0.470496
[126] valid_0's binary_logloss: 0.470461
[127] valid_0's binary_logloss: 0.470453
[128] valid_0's binary_logloss: 0.470476
[129] valid_0's binary_logloss: 0.470482
[130] valid_0's binary_logloss: 0.470472
[131] valid_0's binary_logloss: 0.470463
[132] valid_0's binary_logloss: 0.470432
[133] valid_0's binary_logloss: 0.470437
[134] valid_0's binary_logloss: 0.470412
[135] valid_0's binary_logloss: 0.470408
[136] valid_0's binary_logloss: 0.470363
[137] valid_0's binary_logloss: 0.470332
[138] valid_0's binary_logloss: 0.470325
[139] valid_0's binary_logloss: 0.470301
[140] valid_0's binary_logloss: 0.47029
[141] valid_0's binary_logloss: 0.470277
[142] valid_0's binary_logloss: 0.470267
[143] valid_0's binary_logloss: 0.470269
[144] valid_0's binary_logloss: 0.470284
[145] valid_0's binary_logloss: 0.47028
[146] valid_0's binary_logloss: 0.470289
[147] valid_0's binary_logloss: 0.470284
[148] valid_0's binary_logloss: 0.47028
[149] valid_0's binary_logloss: 0.470262
[150] valid_0's binary_logloss: 0.470252
[151] valid_0's binary_logloss: 0.470263
[152] valid_0's binary_logloss: 0.470245
[153] valid_0's binary_logloss: 0.470234
[154] valid_0's binary_logloss: 0.470231
[155] valid_0's binary_logloss: 0.470219
[156] valid_0's binary_logloss: 0.470187
[157] valid_0's binary_logloss: 0.470198
[158] valid_0's binary_logloss: 0.470208
[159] valid_0's binary_logloss: 0.470194
[160] valid_0's binary_logloss: 0.470202
[161] valid_0's binary_logloss: 0.470212
[162] valid_0's binary_logloss: 0.470182
[163] valid_0's binary_logloss: 0.470189
[164] valid_0's binary_logloss: 0.470199
[165] valid_0's binary_logloss: 0.470197
[166] valid_0's binary_logloss: 0.470195
[167] valid_0's binary_logloss: 0.470171
[168] valid_0's binary_logloss: 0.470141
[169] valid_0's binary_logloss: 0.470143
[170] valid_0's binary_logloss: 0.470142
[171] valid_0's binary_logloss: 0.47016

[172] valid_0's binary_logloss: 0.470152
[173] valid_0's binary_logloss: 0.470118
[174] valid_0's binary_logloss: 0.470107
[175] valid_0's binary_logloss: 0.47011
[176] valid_0's binary_logloss: 0.470103
[177] valid_0's binary_logloss: 0.470108
[178] valid_0's binary_logloss: 0.470115
[179] valid_0's binary_logloss: 0.470092
[180] valid_0's binary_logloss: 0.470081
[181] valid_0's binary_logloss: 0.470085
[182] valid_0's binary_logloss: 0.470089
[183] valid_0's binary_logloss: 0.470088
[184] valid_0's binary_logloss: 0.470071
[185] valid_0's binary_logloss: 0.470057
[186] valid_0's binary_logloss: 0.470062
[187] valid_0's binary_logloss: 0.470058
[188] valid_0's binary_logloss: 0.470057
[189] valid_0's binary_logloss: 0.470063
[190] valid_0's binary_logloss: 0.470061
[191] valid_0's binary_logloss: 0.470069
[192] valid_0's binary_logloss: 0.470061
[193] valid_0's binary_logloss: 0.470053
[194] valid_0's binary_logloss: 0.470045
[195] valid_0's binary_logloss: 0.47005
[196] valid_0's binary_logloss: 0.470052
[197] valid_0's binary_logloss: 0.470063
[198] valid_0's binary_logloss: 0.470061
[199] valid_0's binary_logloss: 0.470056
[200] valid_0's binary_logloss: 0.470027
[201] valid_0's binary_logloss: 0.47003
[202] valid_0's binary_logloss: 0.470007
[203] valid_0's binary_logloss: 0.469976
[204] valid_0's binary_logloss: 0.469943
[205] valid_0's binary_logloss: 0.469951
[206] valid_0's binary_logloss: 0.469941
[207] valid_0's binary_logloss: 0.46995
[208] valid_0's binary_logloss: 0.469962
[209] valid_0's binary_logloss: 0.469961
[210] valid_0's binary_logloss: 0.469949
[211] valid_0's binary_logloss: 0.469952
[212] valid_0's binary_logloss: 0.46996
[213] valid_0's binary_logloss: 0.469955
[214] valid_0's binary_logloss: 0.469957
[215] valid_0's binary_logloss: 0.469945
[216] valid_0's binary_logloss: 0.469948
[217] valid_0's binary_logloss: 0.469945
[218] valid_0's binary_logloss: 0.469943
[219] valid_0's binary_logloss: 0.469931
[220] valid_0's binary_logloss: 0.469931
[221] valid_0's binary_logloss: 0.469918
[222] valid_0's binary_logloss: 0.469912
[223] valid_0's binary_logloss: 0.469909

```

[224] valid_0's binary_logloss: 0.469879
[225] valid_0's binary_logloss: 0.469873
[226] valid_0's binary_logloss: 0.469876
[227] valid_0's binary_logloss: 0.469894
[228] valid_0's binary_logloss: 0.469905
[229] valid_0's binary_logloss: 0.469903
[230] valid_0's binary_logloss: 0.469893
[231] valid_0's binary_logloss: 0.469912
[232] valid_0's binary_logloss: 0.469902
[233] valid_0's binary_logloss: 0.469902
[234] valid_0's binary_logloss: 0.469904
[235] valid_0's binary_logloss: 0.469896
[236] valid_0's binary_logloss: 0.469901
[237] valid_0's binary_logloss: 0.469907
[238] valid_0's binary_logloss: 0.469931
[239] valid_0's binary_logloss: 0.46992
[240] valid_0's binary_logloss: 0.469916
[241] valid_0's binary_logloss: 0.46992
[242] valid_0's binary_logloss: 0.46992
[243] valid_0's binary_logloss: 0.46991
[244] valid_0's binary_logloss: 0.469895
[245] valid_0's binary_logloss: 0.469879
Early stopping, best iteration is:
[225] valid_0's binary_logloss: 0.469873

```

```

LGBMClassifier(boosting_type='gbdt', class_weight=None,
               colsample_bytree=1.0,
                   importance_type='split', learning_rate=0.1, max_depth=-1,
                   metric='binary_logloss', min_child_samples=20,
                   min_child_weight=0.001, min_split_gain=0.0,
n_estimators=100,
                   n_jobs=-1, num_iterations=1000, num_leaves=31,
                   objective='binary', random_state=None, reg_alpha=5,
                   reg_lambda=120, silent=True, subsample=1.0,
                   subsample_for_bin=200000, subsample_freq=0)

```

```

[54] y_pred_test = gbm_clf.predict_proba(x_test)[:, 1]
     y_pred_train = gbm_clf.predict_proba(x_train)[:, 1]
     important_stats(y_train, y_pred_train, "train result summary: ")
     important_stats(y_test, y_pred_test, "test result summary: ")

```

```

-----
train result summary:
recall: 0.7247920665387076
f1_score: 0.770344933104398
accuracy_score: 0.7831516766457454
AUC: 0.8630115360217185
Predicted      0      1    All
True
0              11015   8740  19755
1              11121   8919  20040


```

```

All          22136  17659  39795
-----
-----
test result summary:
recall: 0.7157407407407408
f1_score: 0.7611597374179432
accuracy_score: 0.7757345387302239
AUC: 0.853773792999218
Predicted      0      1   All
True
0             1091    882  1973
1             1135    902  2037
All           2226   1784  4010
-----

```

```
[110] plot_roc_curve(y_train, y_pred_train, y_test, y_pred_test)
```

GridSearch for GBDT

```
[111] from sklearn.model_selection import GridSearchCV
      param_test1 = {'n_estimators':range(20,81,10)}
```

```

gsearch1 = GridSearchCV(estimator =
GradientBoostingClassifier(learning_rate=0.1,
min_samples_split=300,

min_samples_leaf=20,max_depth=8,max_features='sqrt',

subsample=0.8,random_state=10),
                        param_grid = param_test1,
scoring='roc_auc',iid=False,cv=5)
gsearch1.fit(df.drop(columns =
['result','duration']),df['result'])
gsearch1.cv_results_, gsearch1.best_params_, gsearch1.best_score_

({'mean_fit_time': array([1.32538543, 1.89380035, 2.47538829,
2.98179235, 3.5398037 ,
4.12919822, 4.63930659]),
'std_fit_time': array([0.03889608, 0.02579396, 0.02366507, 0.06919284,
0.10178732,
0.06592678, 0.0537543 ]),
'mean_score_time': array([0.04561749, 0.04720511, 0.05919404,
0.06441445, 0.07180109,
0.08361435, 0.08692303]),
'std_score_time': array([0.01460292, 0.00039702, 0.00487201,
0.00100702, 0.00132823,
0.00927824, 0.00326713]),
'param_n_estimators': masked_array(data=[20, 30, 40, 50, 60, 70, 80],
mask=[False, False, False, False, False, False, False],
fill_value='?',
dtype=object),
'params': [{'n_estimators': 20},
{'n_estimators': 30},
{'n_estimators': 40},
{'n_estimators': 50},
{'n_estimators': 60},
{'n_estimators': 70},
{'n_estimators': 80}],
'split0_test_score': array([0.84731079, 0.84979652, 0.85103372,
0.85197477, 0.85263569,
0.85299489, 0.8531087 ]),
'split1_test_score': array([0.85201503, 0.85391403, 0.85478833,
0.85529403, 0.85543664,
0.85559767, 0.85558883]),
'split2_test_score': array([0.84818393, 0.84985972, 0.85031364,
0.85085944, 0.85107263,
0.85120242, 0.85141507]),
'split3_test_score': array([0.8520118 , 0.85381196, 0.85483564,
0.85561528, 0.8556993 ,
0.85586223, 0.85588279]),
'split4_test_score': array([0.84593804, 0.84839916, 0.84974214,
0.85067136, 0.85102791,
0.85125823, 0.8514505 ]),
'mean_test_score': array([0.84909192, 0.85115628, 0.8521427 ,
0.85288298, 0.85317444,
0.85338309, 0.85348918]),
'std_test_score': array([0.00249055, 0.0022711 , 0.00221762,
0.00214893, 0.00204 ,

```

```

0.00202344, 0.00193602]),
'rank_test_score': array([7, 6, 5, 4, 3, 2, 1])),
{'n_estimators': 80},
0.8534891785733354)

```

```

[112] param_test1 = {'n_estimators':range(80,151,10)}
gsearch1 = GridSearchCV(estimator =
GradientBoostingClassifier(learning_rate=0.1,
min_samples_split=300,

min_samples_leaf=20,max_depth=8,max_features='sqrt',

subsample=0.8,random_state=10),
                        param_grid = param_test1,
scoring='roc_auc',iid=False,cv=5)
gsearch1.fit(df.drop(columns =
['result','duration']),df['result'])
gsearch1.cv_results_, gsearch1.best_params_, gsearch1.best_score_

```

```

({'mean_fit_time': array([4.64361835, 5.13062172, 5.66787639,
6.19273343, 6.72101607,
7.27307591, 7.79701047, 8.74241862])),
'std_fit_time': array([0.03414646, 0.11800037, 0.0731966 , 0.12261328,
0.17383439,
0.16389111, 0.07250512, 0.3989696 ]),
'mean_score_time': array([0.08697863, 0.09180603, 0.09779897,
0.10341878, 0.11660056,
0.11560607, 0.12040043, 0.12919893])),
'std_score_time': array([0.00109128, 0.00271011, 0.00074137,
0.00049959, 0.00847716,
0.00233217, 0.00100669, 0.00213085])),
'param_n_estimators': masked_array(data=[80, 90, 100, 110, 120, 130,
140, 150],
mask=[False, False, False, False, False, False, False,
False],
fill_value='?',
dtype=object),
'params': [{'n_estimators': 80},
{'n_estimators': 90},
{'n_estimators': 100},
{'n_estimators': 110},
{'n_estimators': 120},
{'n_estimators': 130},
{'n_estimators': 140},
{'n_estimators': 150}],
'split0_test_score': array([0.8531087 , 0.8530571 , 0.85316144,
0.85306623, 0.85301808,
0.85309912, 0.85309443, 0.85307371]),
'split1_test_score': array([0.85558883, 0.85564495, 0.85574738,
0.85565897, 0.85548331,
0.85535441, 0.85535005, 0.85542428]),
'split2_test_score': array([0.85141507, 0.85155896, 0.85140605,
0.85140464, 0.85153647,

```

```

        0.85153528, 0.85153919, 0.85154953]),
    'split3_test_score': array([0.85588279, 0.85591718, 0.85597176,
0.8558617 , 0.85586745,
        0.85575223, 0.85561992, 0.85554805]),
    'split4_test_score': array([0.8514505 , 0.85148205, 0.85147098,
0.85149088, 0.8516372 ,
        0.85152291, 0.85151866, 0.85149613]),
    'mean_test_score': array([0.85348918, 0.85353205, 0.85355153,
0.85349648, 0.8535085 ,
        0.85345279, 0.85342445, 0.85341834]),
    'std_test_score': array([0.00193602, 0.0019222 , 0.00198813,
0.00194185, 0.00184909,
        0.00181273, 0.00177894, 0.00178132]),
    'rank_test_score': array([5, 2, 1, 4, 3, 6, 7, 8])},
    {'n_estimators': 100},
    0.8535515251856326)

```

```

[114] param_test2 = {'max_depth':range(3,14,2),
    'min_samples_split':range(100,801,200)}
gsearch2 = GridSearchCV(
    estimator = GradientBoostingClassifier(
        learning_rate=0.1,
        n_estimators=120,
        min_samples_leaf=20,
        max_features='sqrt',
        subsample=0.8,
        random_state=10
    ),
    param_grid = param_test2,
    scoring='roc_auc',
    iid=False,
    cv=5)
gsearch2.fit(df.drop(columns =
['result','duration']),df['result'])
gsearch2.cv_results_, gsearch2.best_params_, gsearch2.best_score_

```

```

({'mean_fit_time': array([ 3.45556402,  3.61038938,  3.72701602,
3.67620912,  5.08439488,
        5.28479586,  5.2830235 ,  4.97647386,  6.7170043 ,
6.26970787,
        5.97208714,  5.83681006,  8.37736716,  7.60057206,
7.70422783,
        7.66988192, 10.5684083 ,  9.56854734,  8.91962228,  8.798598
,
        12.91343093, 12.09202356, 10.35147781,  9.29873238])),
    'std_fit_time': array([0.1250714 , 0.05000499, 0.08158398, 0.12801082,
0.08381246,
        0.25093659, 0.13585953, 0.17017597, 0.38197038, 0.16975322,
0.12140917, 0.13991024, 0.43689841, 0.26756902, 0.451388 ,
        0.34377249, 0.16163023, 0.13505017, 0.10145558, 0.16877364,
0.27428557, 0.96775998, 1.02669756, 0.15228346]),
    'mean_score_time': array([0.05162759, 0.06540051, 0.05380282,
0.05259786, 0.08100562,

```

```

0.07601733, 0.08438973, 0.07564855, 0.11180573, 0.10120153,
0.09640412, 0.09418507, 0.13622217, 0.12282095, 0.131393 ,
0.12900147, 0.16540813, 0.16740503, 0.15238347, 0.15100279,
0.19861937, 0.19217443, 0.18200316, 0.15680385]),
'std_score_time': array([0.00301823, 0.02093085, 0.00712394,
0.00632493, 0.00868302,
0.00352076, 0.01046627, 0.00427177, 0.01122829, 0.00685354,
0.00233601, 0.00040692, 0.01814766, 0.00160372, 0.01127262,
0.01121384, 0.01241693, 0.01517542, 0.0059765 , 0.00828008,
0.01509932, 0.02085057, 0.02308627, 0.0021403 ]),
'param_max_depth': masked_array(data=[3, 3, 3, 3, 5, 5, 5, 5, 7, 7, 7,
7, 9, 9, 9, 9, 11, 11,
11, 11, 13, 13, 13, 13],
mask=[False, False, False, False, False, False, False, False,
False,
False, False, False, False, False, False, False, False,
False,
False, False, False, False, False, False, False,
False],
fill_value='?',
dtype=object),
'param_min_samples_split': masked_array(data=[100, 300, 500, 700, 100,
300, 500, 700, 100, 300, 500,
700, 100, 300, 500, 700, 100, 300, 500, 700, 100,
300,
500, 700],
mask=[False, False, False, False, False, False, False, False,
False,
False, False, False, False, False, False, False, False,
False,
False, False, False, False, False, False, False,
False],
fill_value='?',
dtype=object),
'params': [{ 'max_depth': 3, 'min_samples_split': 100},
{ 'max_depth': 3, 'min_samples_split': 300},
{ 'max_depth': 3, 'min_samples_split': 500},
{ 'max_depth': 3, 'min_samples_split': 700},
{ 'max_depth': 5, 'min_samples_split': 100},
{ 'max_depth': 5, 'min_samples_split': 300},
{ 'max_depth': 5, 'min_samples_split': 500},
{ 'max_depth': 5, 'min_samples_split': 700},
{ 'max_depth': 7, 'min_samples_split': 100},
{ 'max_depth': 7, 'min_samples_split': 300},
{ 'max_depth': 7, 'min_samples_split': 500},
{ 'max_depth': 7, 'min_samples_split': 700},
{ 'max_depth': 9, 'min_samples_split': 100},
{ 'max_depth': 9, 'min_samples_split': 300},
{ 'max_depth': 9, 'min_samples_split': 500},
{ 'max_depth': 9, 'min_samples_split': 700},
{ 'max_depth': 11, 'min_samples_split': 100},
{ 'max_depth': 11, 'min_samples_split': 300},
{ 'max_depth': 11, 'min_samples_split': 500},
{ 'max_depth': 11, 'min_samples_split': 700},
{ 'max_depth': 13, 'min_samples_split': 100},
{ 'max_depth': 13, 'min_samples_split': 300},
{ 'max_depth': 13, 'min_samples_split': 500},
{ 'max_depth': 13, 'min_samples_split': 700}],

```

```

'split0_test_score': array([0.84799678, 0.84807903, 0.84810839,
0.84805585, 0.85254784,
    0.85206581, 0.85197321, 0.85224872, 0.85242121, 0.85358496,
    0.85297244, 0.85319115, 0.85197848, 0.85273365, 0.85315618,
    0.85312648, 0.84960534, 0.85186351, 0.85160557, 0.85241442,
    0.84834351, 0.85062906, 0.85112665, 0.85207774]),
'split1_test_score': array([0.85251421, 0.85295653, 0.85256148,
0.85272752, 0.85486332,
    0.85520937, 0.85494075, 0.85493413, 0.8559036 , 0.85593558,
    0.85578692, 0.85605588, 0.85459506, 0.85525841, 0.85522146,
    0.85559023, 0.85328955, 0.85378909, 0.85512607, 0.85520454,
    0.85138227, 0.85367802, 0.85426616, 0.85504019]),
'split2_test_score': array([0.84742484, 0.84802475, 0.84804271,
0.84769015, 0.85133416,
    0.85150392, 0.85129539, 0.85110923, 0.85137496, 0.85176178,
    0.8515078 , 0.85209579, 0.85083211, 0.85141843, 0.85187714,
    0.85186189, 0.84955372, 0.85036691, 0.85102396, 0.85224837,
    0.84757987, 0.84947853, 0.85057135, 0.85122504]),
'split3_test_score': array([0.85118431, 0.85156696, 0.85157943,
0.85167321, 0.85567397,
    0.85527397, 0.855096 , 0.85527526, 0.855685 , 0.85551696,
    0.85575569, 0.85565237, 0.85494414, 0.85536433, 0.85599175,
    0.85594658, 0.8528644 , 0.85497016, 0.85605099, 0.85545646,
    0.85045409, 0.85392792, 0.85460427, 0.85501261]),
'split4_test_score': array([0.84557792, 0.84519506, 0.84560721,
0.84561366, 0.85006306,
    0.85052858, 0.85039992, 0.85019692, 0.85095416, 0.85171547,
    0.85110735, 0.85148942, 0.84951247, 0.85175528, 0.8513141 ,
    0.85206876, 0.84893426, 0.85021242, 0.85076372, 0.85114496,
    0.8469133 , 0.84990433, 0.85041393, 0.85108742]),
'mean_test_score': array([0.84893961, 0.84916447, 0.84917984,
0.84915208, 0.85289647,
    0.85291633, 0.85274105, 0.85275285, 0.85326779, 0.85370295,
    0.85342604, 0.85369692, 0.85237245, 0.85330602, 0.85351213,
    0.85371879, 0.85084945, 0.85224042, 0.85291406, 0.85329375,
    0.84893461, 0.85152357, 0.85219647, 0.8528886 ]),
'std_test_score': array([0.00254169, 0.00277074, 0.00254552,
0.00264572, 0.00210588,
    0.00196144, 0.00192587, 0.00203022, 0.00211862, 0.00178927,
    0.0020131 , 0.00184829, 0.00211002, 0.00169375, 0.00182765,
    0.00173129, 0.00183896, 0.00187695, 0.00222 , 0.00172115,
    0.00170706, 0.00189878, 0.0018463 , 0.0017782 ]),
'rank_test_score': array([23, 21, 20, 22, 11, 9, 14, 13, 8, 2, 5,
3, 15, 6, 4, 1, 19,
    16, 10, 7, 24, 18, 17, 12])),
{'max_depth': 9, 'min_samples_split': 700},
0.8537187869015375)

```

```

[116] param_test3 = {'min_samples_leaf':range(60,101,10)}
gsearch3 = GridSearchCV(
    estimator = GradientBoostingClassifier(
        learning_rate=0.1,
        n_estimators=120,
        max_depth=7,
        min_samples_split=700,

```

```

        max_features='sqrt',
        subsample=0.8,
        random_state=10
    ),
    param_grid = param_test3,
    scoring='roc_auc',
    iid=False,
    verbose=1,
    cv=5
)

gsearch3.fit(df.drop(columns =
['result', 'duration']), df['result'])
gsearch3.cv_results_, gsearch3.best_params_, gsearch3.best_score_

```

Fitting 5 folds for each of 5 candidates, totalling 25 fits
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[Parallel(n_jobs=1)]: Done 25 out of 25 | elapsed: 2.7min finished

```

({'mean_fit_time': array([6.06449895, 6.34776134, 6.46800838,
6.06420708, 6.41620145]),
 'std_fit_time': array([0.1295597 , 0.08398248, 0.31314299, 0.35224719,
0.3227051 ]),
 'mean_score_time': array([0.10140619, 0.10560927, 0.11020746,
0.10400147, 0.10178876]),
 'std_score_time': array([0.01030432, 0.0075005 , 0.01550756,
0.00918721, 0.00205163]),
 'param_min_samples_leaf': masked_array(data=[60, 70, 80, 90, 100],
      mask=[False, False, False, False, False],
      fill_value='?',
      dtype=object),
 'params': [{'min_samples_leaf': 60},
 {'min_samples_leaf': 70},
 {'min_samples_leaf': 80},
 {'min_samples_leaf': 90},
 {'min_samples_leaf': 100}],
 'split0_test_score': array([0.85278824, 0.85324258, 0.85328524,
0.85319134, 0.85283255]),
 'split1_test_score': array([0.85619283, 0.85620537, 0.85633786,
0.85643664, 0.85641502]),
 'split2_test_score': array([0.85210429, 0.85168207, 0.85245235,
0.8522972 , 0.85226799]),
 'split3_test_score': array([0.85606564, 0.85636268, 0.85600638,
0.85561514, 0.85574515]),
 'split4_test_score': array([0.85183398, 0.85146712, 0.8512513 ,
0.85090273, 0.85107601]),
 'mean_test_score': array([0.853797 , 0.85379196, 0.85386663,
0.85368861, 0.85366734]),
 'std_test_score': array([0.00192992, 0.00212563, 0.00199317,
0.00205949, 0.00206093]),
 'rank_test_score': array([2, 3, 1, 4, 5])},
 {'min_samples_leaf': 80},
 0.8538666256737037)

```

```
[16] gbd_t_best = GradientBoostingClassifier(
    learning_rate=0.1,
    n_estimators=100,
    max_depth=9,
    min_samples_leaf =80,
    min_samples_split =700,
    max_features='sqrt',
    subsample=0.8,
    random_state=10
)
gbd_t_best.fit(df.drop(columns =
['result','duration']),df['result'])
```

```
GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse',
init=None,
                                learning_rate=0.1, loss='deviance',
max_depth=9,
                                max_features='sqrt', max_leaf_nodes=None,
                                min_impurity_decrease=0.0,
min_impurity_split=None,
                                min_samples_leaf=80, min_samples_split=700,
                                min_weight_fraction_leaf=0.0,
n_estimators=100,
                                n_iter_no_change=None, presort='deprecated',
                                random_state=10, subsample=0.8, tol=0.0001,
                                validation_fraction=0.1, verbose=0,
                                warm_start=False)
```


```
[17] y_pred_test = gbd_t_best.predict_proba(x_test)[: , 1]
y_pred_train = gbd_t_best.predict_proba(x_train)[: , 1]
important_stats(y_train, y_pred_train, "train result summary: ")
important_stats(y_test, y_pred_test, "test result summary: ")
```

```
-----
train result summary:
recall: 0.7176787777635126
f1_score: 0.7702642185400805
accuracy_score: 0.7859000288933834
AUC: 0.8671379869262116
Predicted      0      1    All
True
0              11315   8599  19914
1              11361   8536  19897
All            22676  17135  39811
-----
-----
test result summary:
recall: 0.7218407297324998
f1_score: 0.7745518530737929
accuracy_score: 0.7894088037392778
```


AUC: 0.871171321641422

Predicted	0	1	All
True			
0	1128	867	1995
1	1127	888	2015
All	2255	1755	4010

```
[18] plot_roc_curve(y_train, y_pred_train, y_test, y_pred_test)
```




Features importance

```
[19] gbdt_importance = gbdt_best.feature_importances_
```

```
[122] from matplotlib.pyplot import figure
      figure(num=None, figsize = (10,10))
      indices = np.argsort(gbdt_importance)
      plt.figure(1)
      plt.title('GBDT Classifier Feature Importance')
      plt.barh(range(len(indices)), gbdt_importance[indices], color =
        'b', align = 'center')
      gbdt_feat_names = x_train.columns
      plt.yticks(range(len(indices)), gbdt_feat_names[indices])
      plt.xlabel('Relative Importance: Gini')
```

```
Text(0.5, 0, 'Relative Importance: Gini')
```



Logistic Regression for predicting probabilities

```
[79] from sklearn.linear_model import LogisticRegressionCV
lr = LogisticRegressionCV(solver = 'saga',
                           penalty = 'elasticnet',
```

```

        l1_ratios = [0.1, 0.2, 0.3],
        Cs = 20,
        n_jobs = -1,
        random_state = 0,
        class_weight = 0.9
    )
    lr.fit(x_train,y_train)

```

D:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection_split.py:1978: FutureWarning: The default value of cv will change from 3 to 5 in version 0.22. Specify it explicitly to silence this warning.

```
warnings.warn(CV_WARNING, FutureWarning)
```

D:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
"the coef_ did not converge", ConvergenceWarning)

D:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
"the coef_ did not converge", ConvergenceWarning)

D:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
"the coef_ did not converge", ConvergenceWarning)

D:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
"the coef_ did not converge", ConvergenceWarning)

D:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
"the coef_ did not converge", ConvergenceWarning)

D:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
"the coef_ did not converge", ConvergenceWarning)

D:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
"the coef_ did not converge", ConvergenceWarning)

D:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
"the coef_ did not converge", ConvergenceWarning)

D:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
"the coef_ did not converge", ConvergenceWarning)

D:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge

```

    "the coef_ did not converge", ConvergenceWarning)
D:\ProgramData\Anaconda3\lib\site-
packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The
max_iter was reached which means the coef_ did not converge
    "the coef_ did not converge", ConvergenceWarning)
D:\ProgramData\Anaconda3\lib\site-
packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The
max_iter was reached which means the coef_ did not converge
    "the coef_ did not converge", ConvergenceWarning)
D:\ProgramData\Anaconda3\lib\site-
packages\sklearn\linear_model\sag.py:337: ConvergenceWarning: The
max_iter was reached which means the coef_ did not converge
    "the coef_ did not converge", ConvergenceWarning)

LogisticRegressionCV(Cs=20, class_weight=0.9, cv='warn', dual=False,
                      fit_intercept=True, intercept_scaling=1.0,
                      l1_ratios=[0.1, 0.2, 0.3], max_iter=100,
                      multi_class='warn', n_jobs=-1,
penalty='elasticnet',
                      random_state=0, refit=True, scoring=None,
solver='saga',
                      tol=0.0001, verbose=0)

```

```

[80] y_pred_test = lr.predict_proba(x_test)[: , 1]
y_pred_train = lr.predict_proba(x_train)[: , 1]
important_stats(y_train, y_pred_train, "train result summary: ")
important_stats(y_test, y_pred_test, "test result summary: ")

```

```

-----
train result summary:
recall: 0.5398912348048625
f1_score: 0.6377824805381301
accuracy_score: 0.6922805271521904
AUC: 0.7677302786258187
Predicted      0      1    All
True
0              12899   6856  19755
1              13041   6999  20040
All            25940  13855  39795
-----

test result summary:
recall: 0.5417695473251029
f1_score: 0.6383030303030303
accuracy_score: 0.6934456544072324
AUC: 0.7677087607714403
Predicted      0      1    All
True
0              1280   693  1973
1              1292   745  2037

```

All 2572 1438 4010

```
[81] plot_roc_curve(y_train, y_pred_train, y_test, y_pred_test)
```

